

Methods Overview

Methods Overview I

The pipeline is a sequence of deterministic stages, each reading the previous stage's committed artifacts and writing its own. Business logic lives in tested `src/` modules; the numbered `scripts/` are thin orchestrators that wire I/O, configuration loading, logging, and stage sequencing. The architecture follows the thin orchestrator pattern: no computational logic resides in scripts.

Pipeline Stages I

1. **Retrieval** (`01_literature_search.py`) — dispatch the configured query across 10 engines (arXiv, OpenAlex, Semantic Scholar, Crossref, PubMed, SovietRxiv, ChinaRxiv, Europe PMC, bioRxiv, and medRxiv), merge, and de-duplicate into `corpus.jsonl`. Each engine is an isolated adapter exposing a uniform `search(query) -> list[Paper]` interface; engines that are keyless need no credentials, while Semantic Scholar uses a key when present. SovietRxiv and ChinaRxiv share a unified API with an optional `X-API-Email` header for the polite rate-limit pool (300/min vs 30/min anonymous).

Pipeline Stages II

2. **Meta-analysis** (02_meta_analysis_pipeline.py) — subfield classification, temporal metrics, TF-IDF, non-negative matrix factorization topics, and the citation network. This stage reads `corpus.jsonl` and emits `subfield_classification.json`, `temporal_analysis.json`, `tfidf_data.json`, `topics.json`, `citation_network.json`, and `citation_graph.gml`.
3. **Knowledge graph** (03_build_knowledge_graph.py, optional/LLM-gated) — extract assertions and score the 6 configured hypotheses. Outputs `nanopublications.jsonl`, `hypothesis_scores.json`, and `assertion_summary.json`.
4. **Figures** (04_generate_figures.py) — render 21 publication-ready visualizations from the analysis JSON outputs. All figures use a colourblind-safe palette (Wong 2011), high-contrast labels at ≥ 16 pt, and a headless matplotlib backend (Agg).

Pipeline Stages III

5. **Injection** (`05_inject_variables.py`) — compute manuscript variables from the artifacts above and substitute them into these Markdown sections. An unresolved placeholder is a hard error, not a silent gap.
6. **Fulltext assessment** (`06_fulltext_assessment.py`) — report abstract coverage (61.6%), open-access status (24.6%), and PDF availability (54.6%) across the corpus.

Reproducibility Model I

The system runs **offline and deterministically** by default: a committed synthetic seed corpus drives every stage with fixed seeds (seed = 42 for NMF, SVD, and graph layouts), so re-running produces byte-identical outputs. A live run with engines enabled and credentials supplied replaces the seed corpus with real records — as in this instance, which retrieved 2334 live records. The template is domain-agnostic: the search term, query, keyword set, subfield taxonomy, and hypotheses all come from `manuscript/config.yaml`.

Configuration Surface I

A single manuscript/config.yaml controls:

- ▶ **Search parameters:** term, query string, per-engine queries, relevance keywords, start year, max results, resume/clear behaviour
- ▶ **Engine toggles:** arXiv, OpenAlex, Semantic Scholar, Crossref, PubMed, SovietRxiv, ChinaRxiv, Europe PMC, bioRxiv, and medRxiv (each independently enabled or disabled)
- ▶ **SovietRxiv/ChinaRxiv settings:** optional api_email for the polite pool, source filter (russiарxiv or chinaxiv)
- ▶ **Full-text download:** opt-in Unpaywall resolution with unpaywall_email
- ▶ **Embeddings:** method (tfidf_svd or transformer), dimensionality, max features
- ▶ **Knowledge graph:** checkpoint interval, LLM model, base URL, temperature, max tokens

Configuration Surface II

- ▶ **Hypothesis definitions:** 6 named hypotheses with scope labels
- ▶ **Subfield taxonomy:** 6 buckets, each with a keyword list
- ▶ **Paper metadata:** title, authors, DOI, keywords, license, repository URL