

Results: Language, Topics, and Embeddings

## RQ3: Topical and Linguistic Structure I

Text analysis operates over titles, abstracts, and (when available) full text. A TF-IDF representation over a 500-feature vocabulary feeds non-negative matrix factorization, which extracts 5 latent topics cross-cutting the subfield taxonomy. The top vocabulary terms are: modafinil, treatment, study, results, patients, effects, sleep, used, clinical, use, drug, studies, mg, using, sleepiness, disorder, associated, cognitive, narcolepsy, significant.

**Table 3. NMF topics extracted from the corpus.**

| Topic | Top terms                                                                        |
|-------|----------------------------------------------------------------------------------|
| 0     | modafinil, mg, kg, effects, rats, mice, induced, sleep                           |
| 1     | fatigue, patients, placebo, modafinil, scale, armodafinil, treatment, depression |
| 2     | h4, methods, results, conclusion, 95, analysis, ci, risk                         |
| 3     | use, cognitive, drugs, adhd, studies, drug, cocaine, methylphenidate             |
| 4     | sleep, narcolepsy, sleepiness, cataplexy, daytime, patients, eds, excessive      |

## RQ3: Topical and Linguistic Structure II

The topics reveal the thematic structure of the literature: Topic 0 centres on cognitive enhancement and neuroenhancement; Topic 1 addresses ADHD treatment and clinical evidence; Topic 2 covers pharmacological dose-response studies (including animal models); Topic 3 focuses on sleep disorders (narcolepsy, excessive daytime sleepiness); and Topic 4 addresses fatigue in psychiatric populations. These topics cross-cut the keyword-based subfield taxonomy, revealing connections that the explicit classification does not capture.

# RQ3: Topical and Linguistic Structure III

NMF Topic — Top Terms





Figure 1: Topic-term bar charts for Modafinil. Each panel shows the top weighted terms for one of 5 NMF topics, with bar length proportional to the topic-term weight in the  $\mathbf{H}$  matrix.

# Document Embeddings I

Offline deterministic embeddings (TF-IDF followed by truncated SVD) place every document in a shared 50-dimensional vector space. Embedding the same text twice yields identical vectors, so the derived similarity matrix, nearest-neighbour lists, clusters, and two-dimensional projection are all reproducible.


# Document Embeddings II



# Document Embeddings III

Figure 2: PCA projection of document embeddings for Modafinil. Each point represents one document projected onto the first two principal components of the TF-IDF/SVD embedding. Colours indicate subfield assignment, showing how the topical geography relates to the keyword taxonomy.

# Document Embeddings IV






# Document Embeddings V

Figure 3: Hierarchical clustering dendrogram of document embeddings. The tree shows the similarity structure of the corpus: documents that join low in the tree are semantically similar, while high-level splits separate the major topical clusters.

# Term Analysis I

The TF-IDF term heatmap reveals which terms discriminate between subfields: terms with high between-subfield variance (rather than high global mean) are selected for display.



## Term Analysis II

Figure 4: Term heatmap for Modafinil. Each cell shows the mean TF-IDF weight of a term within a subfield. Terms are selected by between-subfield variance to highlight discriminative vocabulary rather than globally frequent terms.

# Named Entity Analysis I

Named entity extraction over the 1437 abstracts identified 30 unique entities. The most frequent entities reflect the clinical and pharmacological vocabulary of the modafinil literature.

**Table 4. Top named entities in abstracts.**

| Entity      | Frequency |
|-------------|-----------|
| ADHD        | 392       |
| CI          | 343       |
| EDS         | 326       |
| OSA         | 272       |
| MOD         | 181       |
| ESS         | 148       |
| RESULTS     | 148       |
| MS          | 145       |
| DAT         | 134       |
| NT1         | 131       |
| SD          | 130       |
| CONCLUSIONS | 129       |
| MD          | 102       |
| CE          | 101       |
| IH          | 101       |

# Named Entity Analysis II




Figure 5: Top named entities for Modafinil. The horizontal bar chart shows the 20 most frequently extracted named entities from abstracts, revealing the dominant drugs, conditions, and concepts in the literature.

# Named Entity Analysis III

**Table 5. Top keyphrases by TF-IDF score.**

| Keyphrase          | Score  |
|--------------------|--------|
| available          | 0.3333 |
| abstract           | 0.3333 |
| abstract available | 0.3333 |
| content            | 0.1053 |
| access             | 0.1053 |
| sub                | 0.0881 |
| md                 | 0.0870 |
| jama               | 0.0826 |
| cleveland          | 0.0763 |
| modafinil          | 0.0741 |
| depression         | 0.0741 |
| cognitive          | 0.0714 |
| substance          | 0.0694 |
| drug               | 0.0667 |
| conditions         | 0.0667 |

# Embedding Similarity and Clustering I

The TF-IDF/SVD embeddings place every document in a 50-dimensional vector space. K-means clustering with  $k = 5$  clusters partitions the corpus into topically coherent groups. The top similar document pairs, ranked by cosine similarity, reveal the most closely related works in the corpus.

**Table 6. Top 10 most similar document pairs.**

| Paper A                        | Paper B                        | Similarity |
|--------------------------------|--------------------------------|------------|
| doi:10.1176/appi.ajp.163.12.21 | doi:10.1176/ajp.2006.163.12.21 | 1.0000     |
| doi:10.1176/ajp.2006.163.12.21 | doi:10.1176/appi.ajp.163.12.21 | 1.0000     |
| doi:10.1345/aph.1h302          | doi:10.1136/bcr.08.2011.4652   | 0.9660     |
| doi:10.4088/jcp.09m05900gry    | doi:10.1186/s40345-015-0034-0  | 0.9566     |
| doi:10.1111/bdi.12859          | doi:10.1111/acps.12712         | 0.9558     |
| doi:10.1111/j.1360-0443.2008.0 | doi:10.1111/j.1465-3362.2012.0 | 0.9536     |
| doi:10.1513/annalsats.202006-6 | doi:10.3760/cma.j.cn112147-202 | 0.9532     |
| doi:10.1164/ajrccm.164.9.21030 | doi:10.1093/sleep/28.4.464     | 0.9513     |
| doi:10.1093/sleep/28.4.464     | doi:10.1164/ajrccm.164.9.21030 | 0.9513     |
| doi:10.1002/cncr.25083         | doi:10.1200/jco.2005.23.16_sup | 0.9507     |

# Embedding Similarity and Clustering II

## Top 15 Most Similar Document Pairs





Figure 6: Document similarity for Modafinil. The horizontal bar chart shows the 15 most similar document pairs ranked by cosine similarity of their TF-IDF/SVD embeddings. High-similarity pairs share topical and lexical content.



# Embedding Similarity and Clustering III

## Literature Corpus — Term Cloud



## Embedding Similarity and Clustering IV

Figure 7: Term cloud for Modafinil. Term sizes are proportional to their TF-IDF weights across the corpus, providing a visual summary of the dominant vocabulary.

# Embedding Similarity and Clustering V

## Term Co-occurrence Matrix

